Hewlett Packard Enterprise
Data Science Institute
UNIVERSITY OF HOUSTON

311 Introduction to Machine Learning

Summer 2024

Instructor: Ioannis Konstantinidis

Overview

- Metrics and Scoring:
  - Confusion Matrix
  - Error Functions
  - Regularization
- Hands-on examples
  - Classification
  - Regression

Developed by Ioannis Konstantinidis

Hewlett Packard Enterprise Data Science Institute
UNIVERSITY OF HOUSTON

Ways of Quantifying the Predictive Capability of Classification Tasks

Confusion Matrix

Confusion Matrix

Columns: predictions made by the classifier (labels y)

Rows: actual observations (points X)

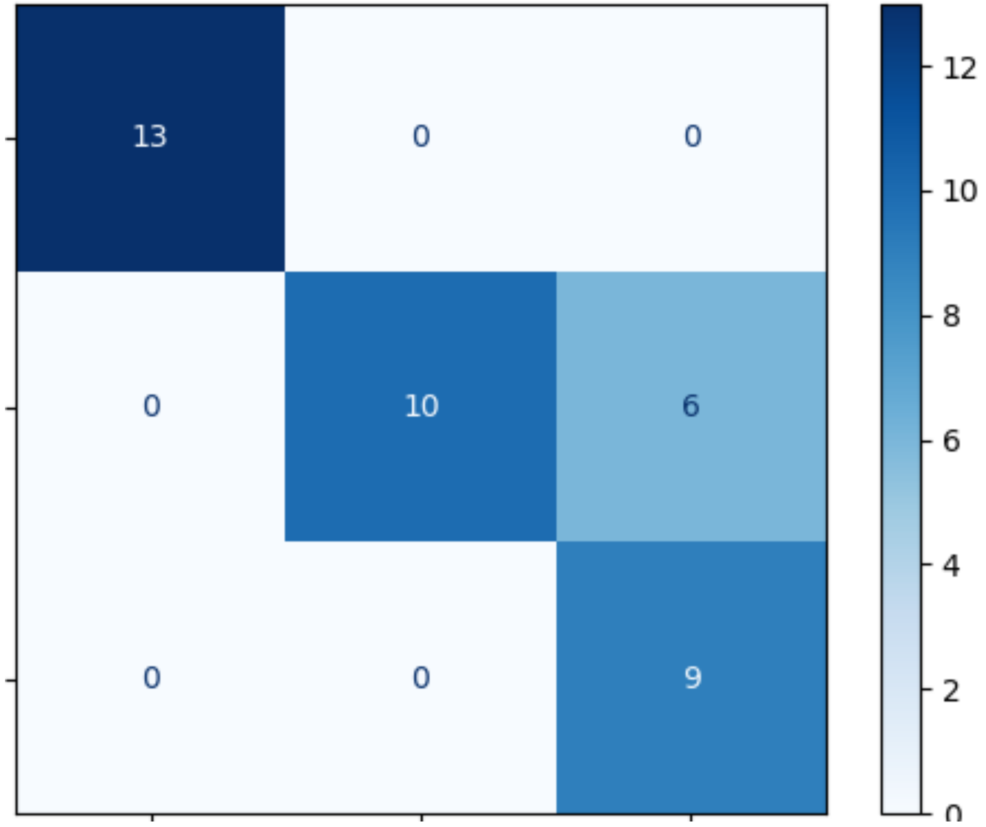

Confusion Matrix

Columns: predictions made by the classifier (labels y)

Rows: actual observations (points X)

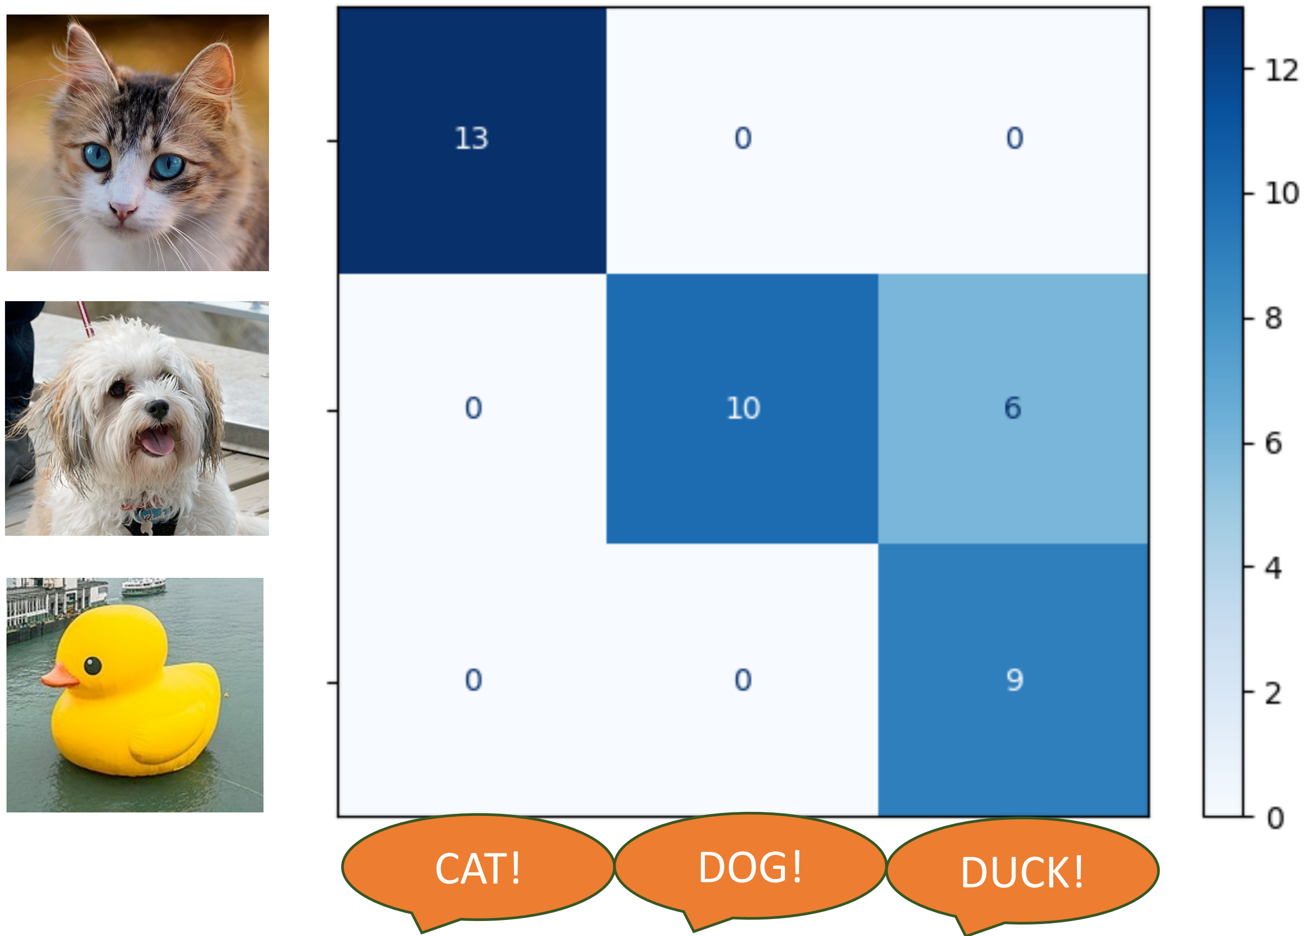

Confusion Matrix

- Diagonal: # of points for which predicted label = true label
- Off-diagonal: # of points that are mislabeled by the classifier
- The smaller the off-diagonal values of the confusion matrix, the better

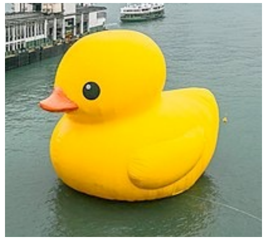

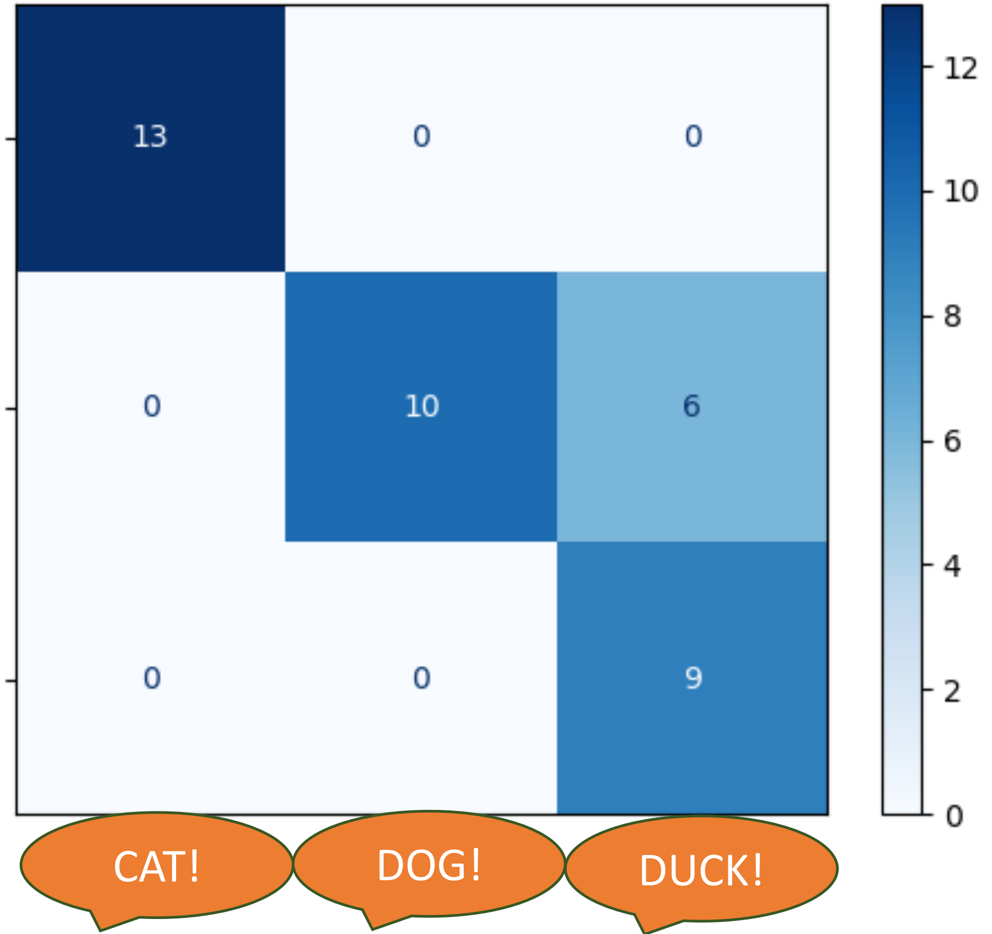

Confusion Matrix

$$\textbf{Accuracy} = \frac{\text{Diagonal}}{\text{All}}$$

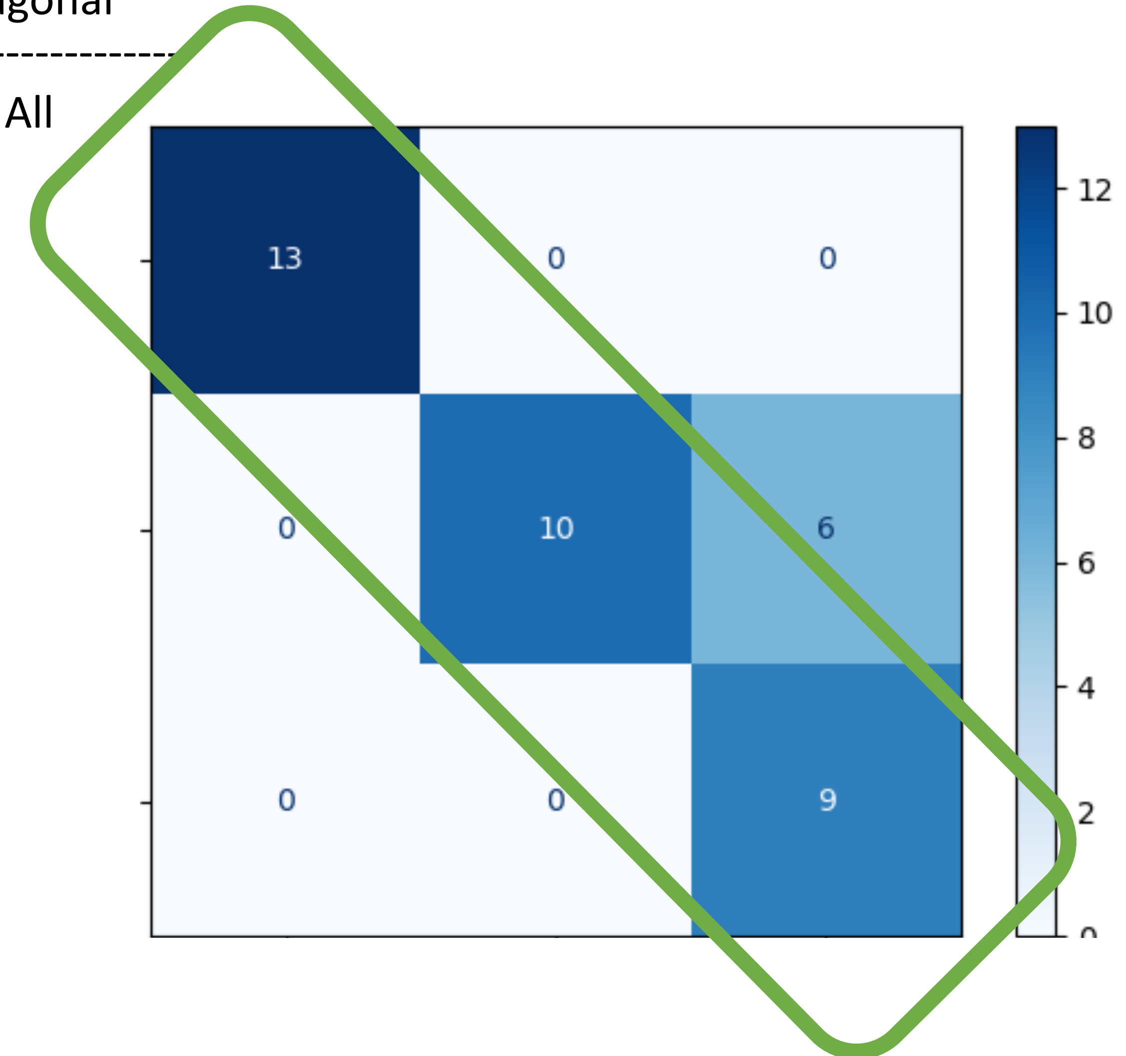

Confusion Matrix

Only one type of error:
all the off-diagonal entries

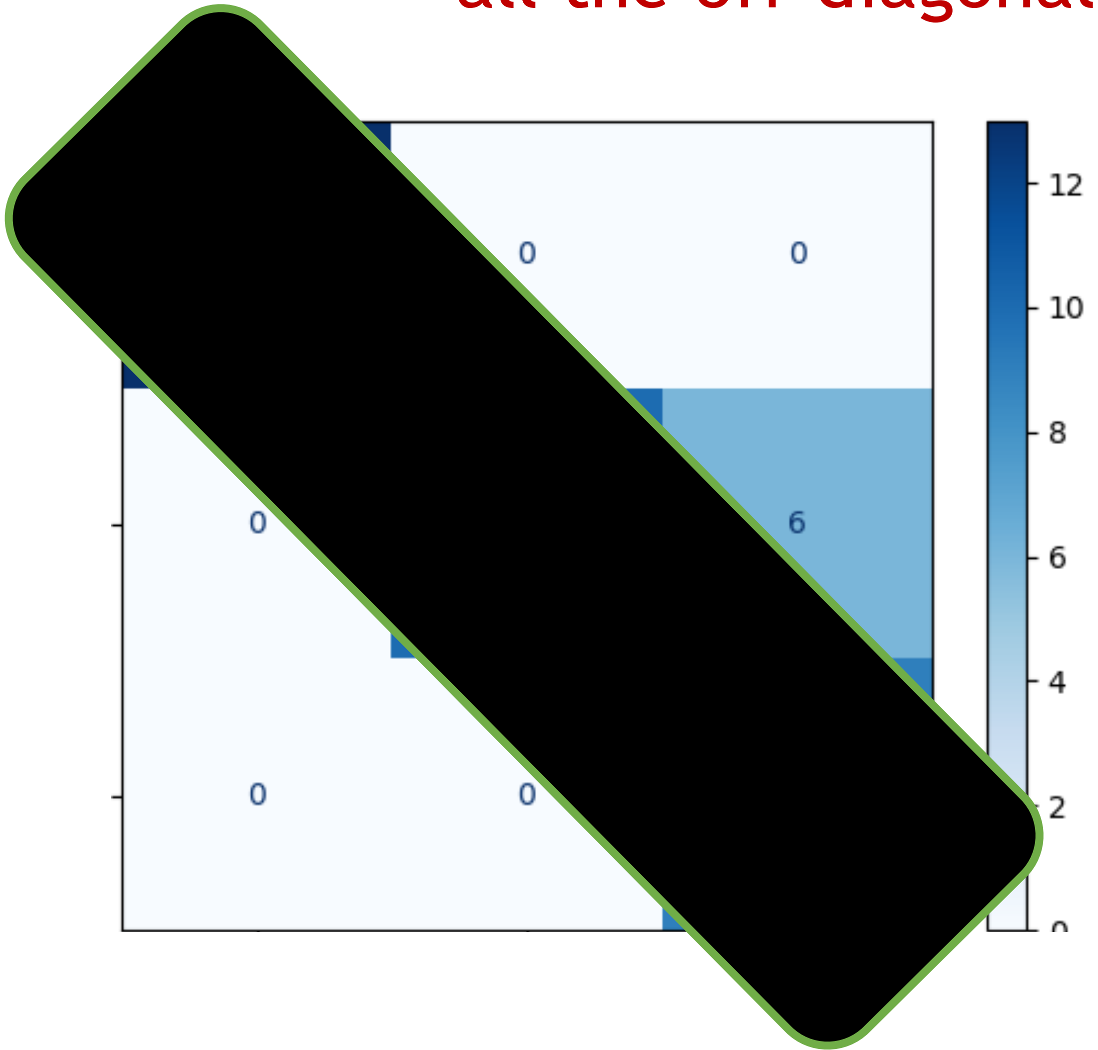

Focus on a single label

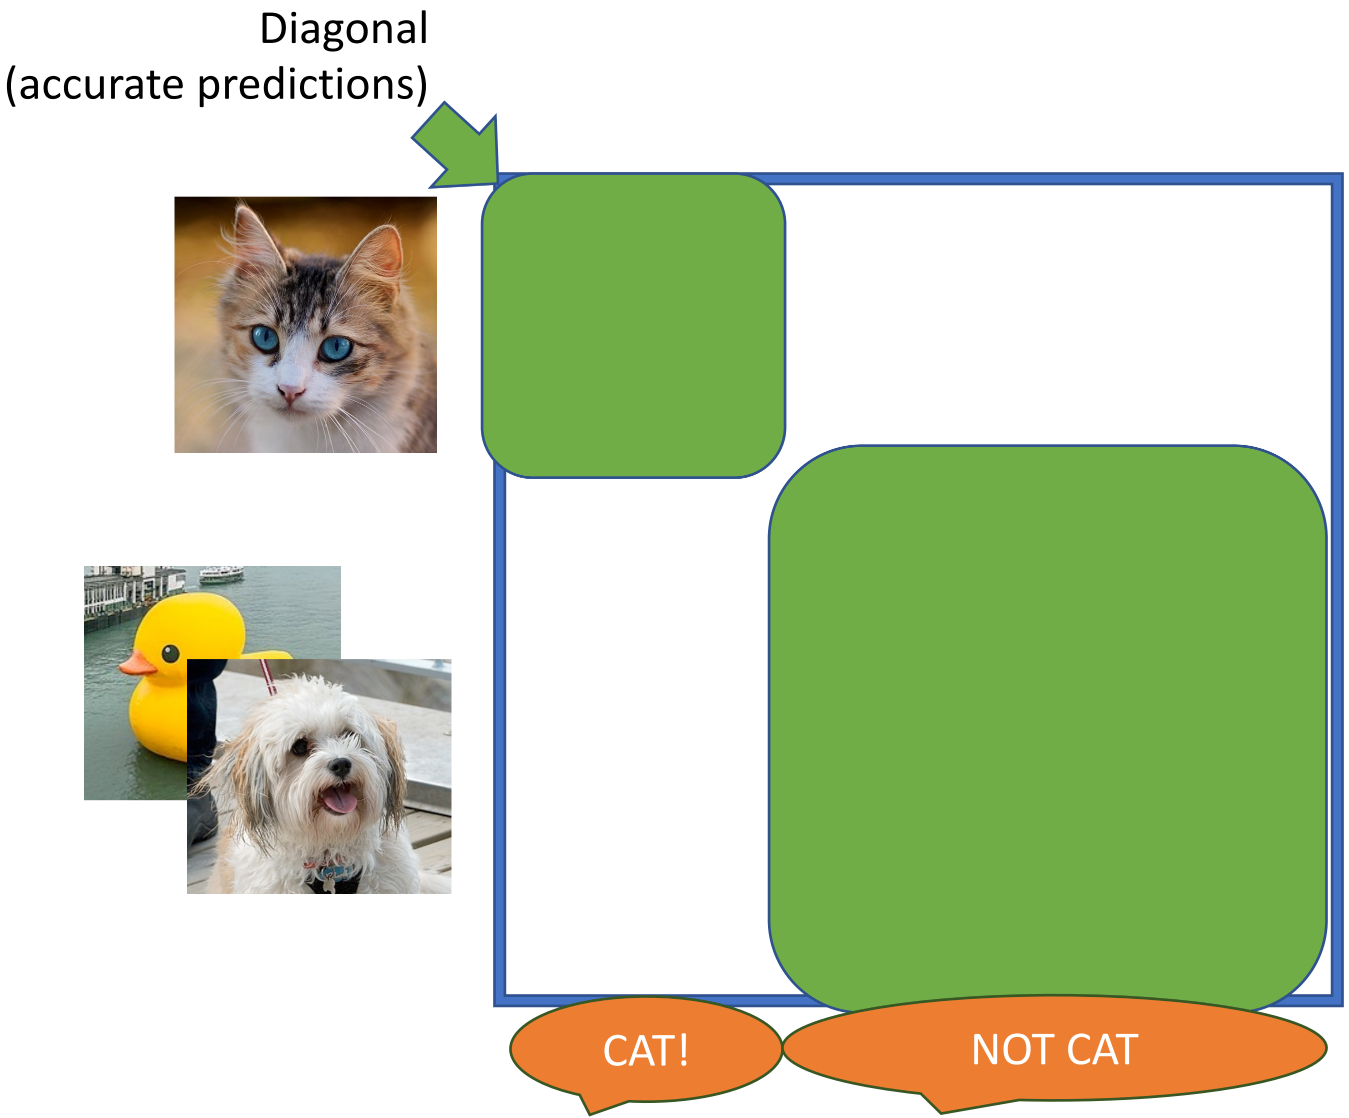

Focus on a single label

Off-diagonal
(errors)

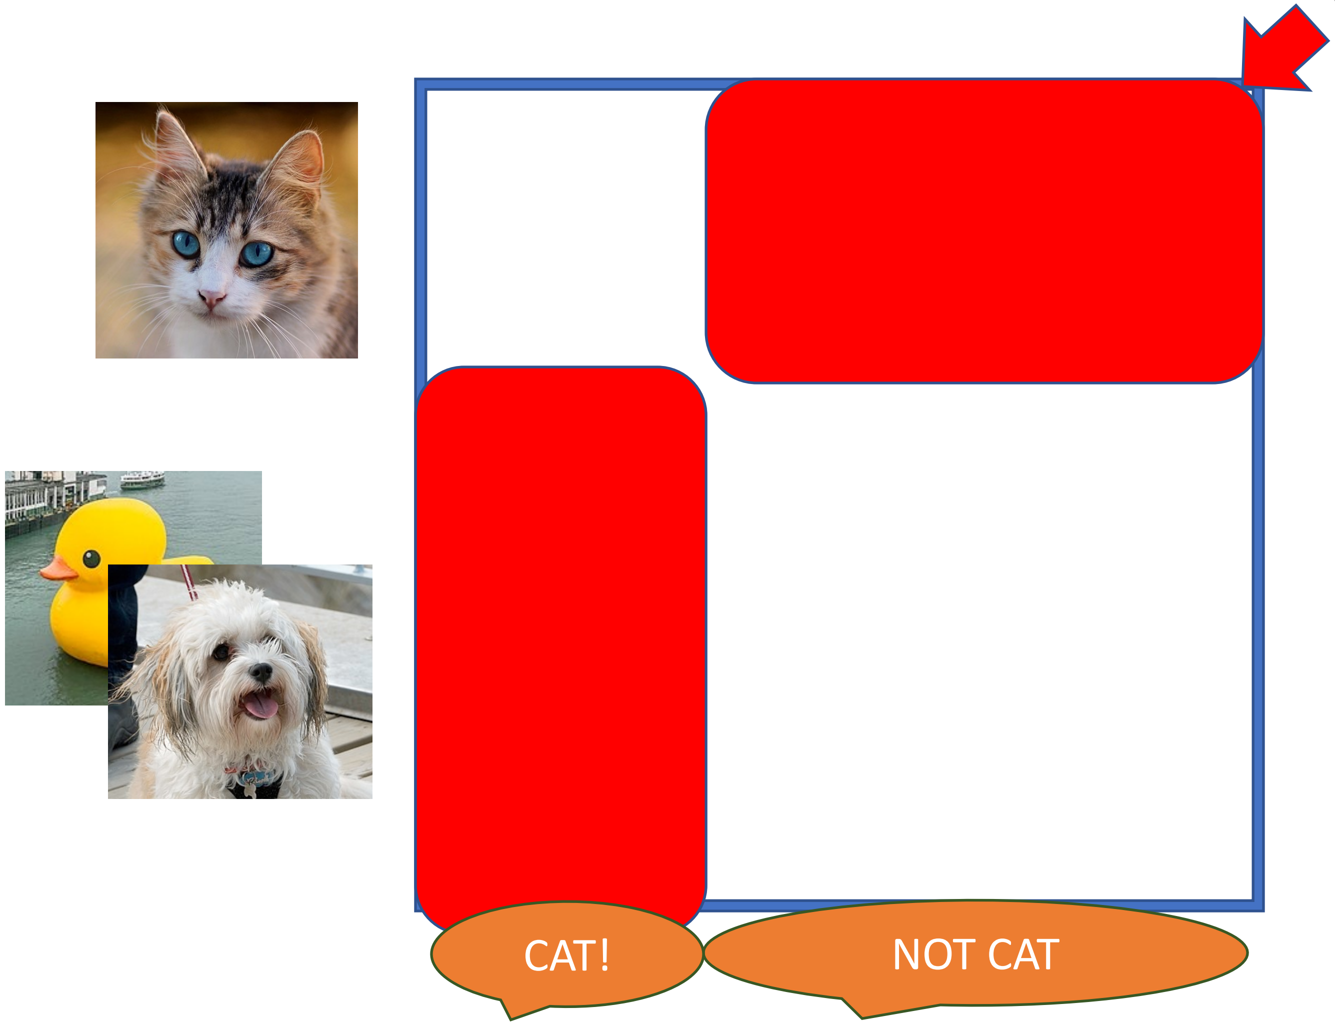

Focus on a single label

TWO types of error:

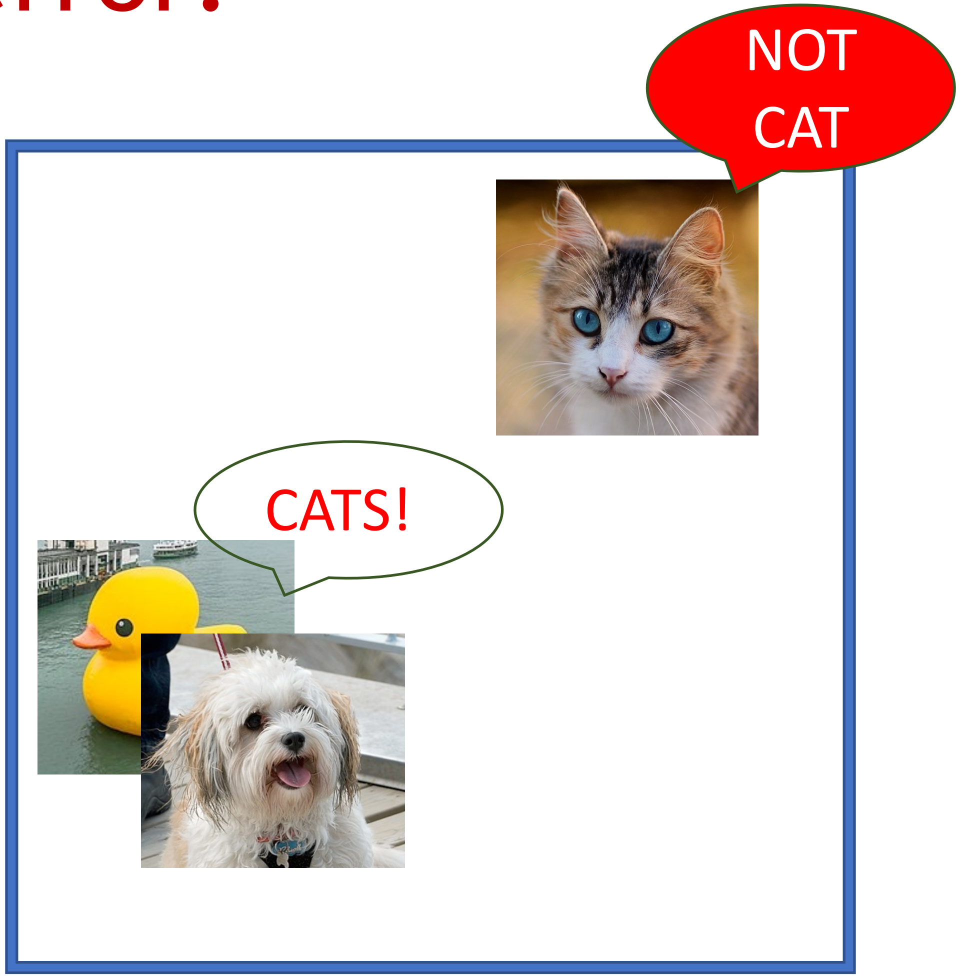

Focus on a single label

TWO types of error:

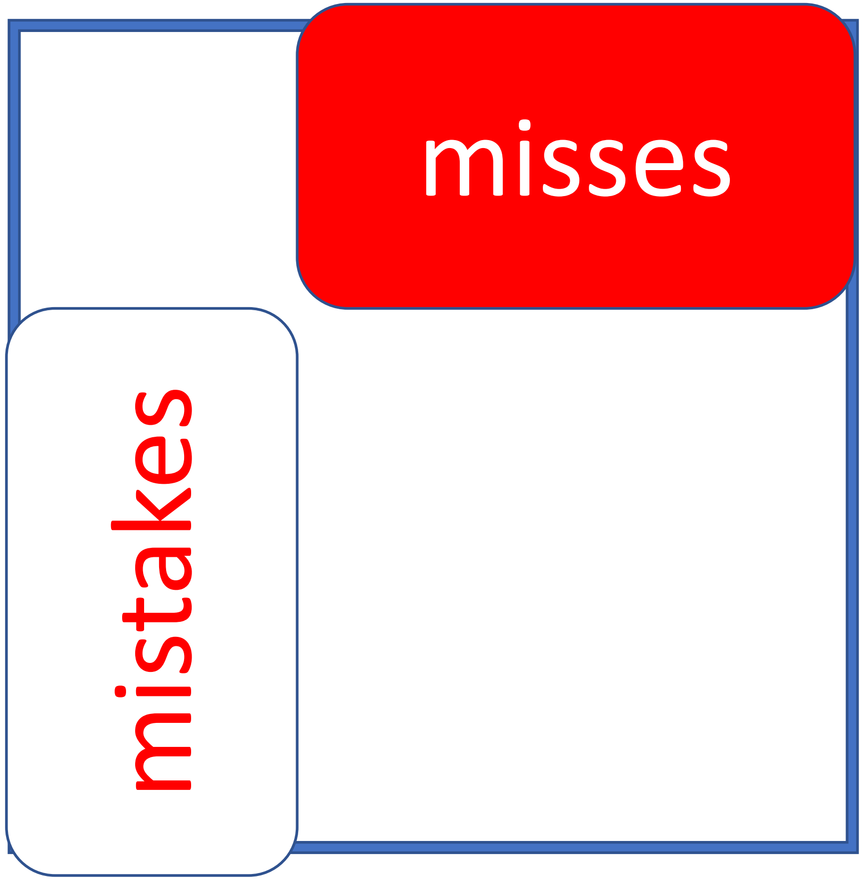

Focus on a single label

TWO types of error and two correct types

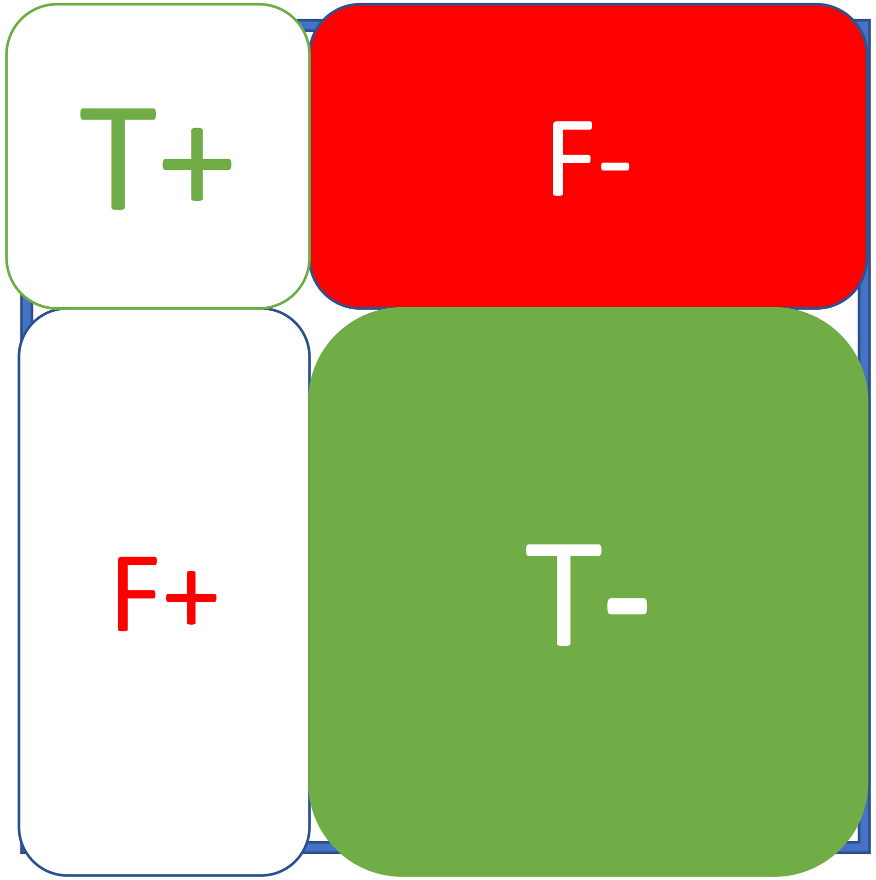

Focus on a single label

$$\textbf{Precision} = \frac{T+}{T+ \; + \; F+}$$

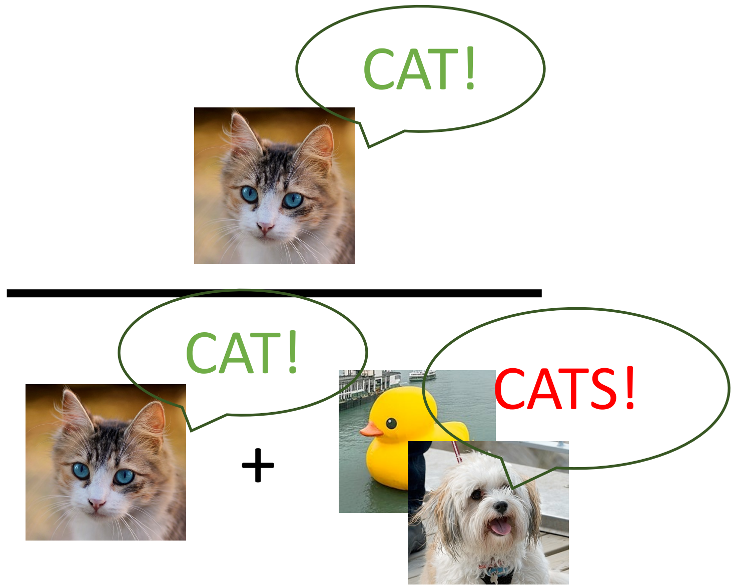

Positives column
(predicted to be the target)

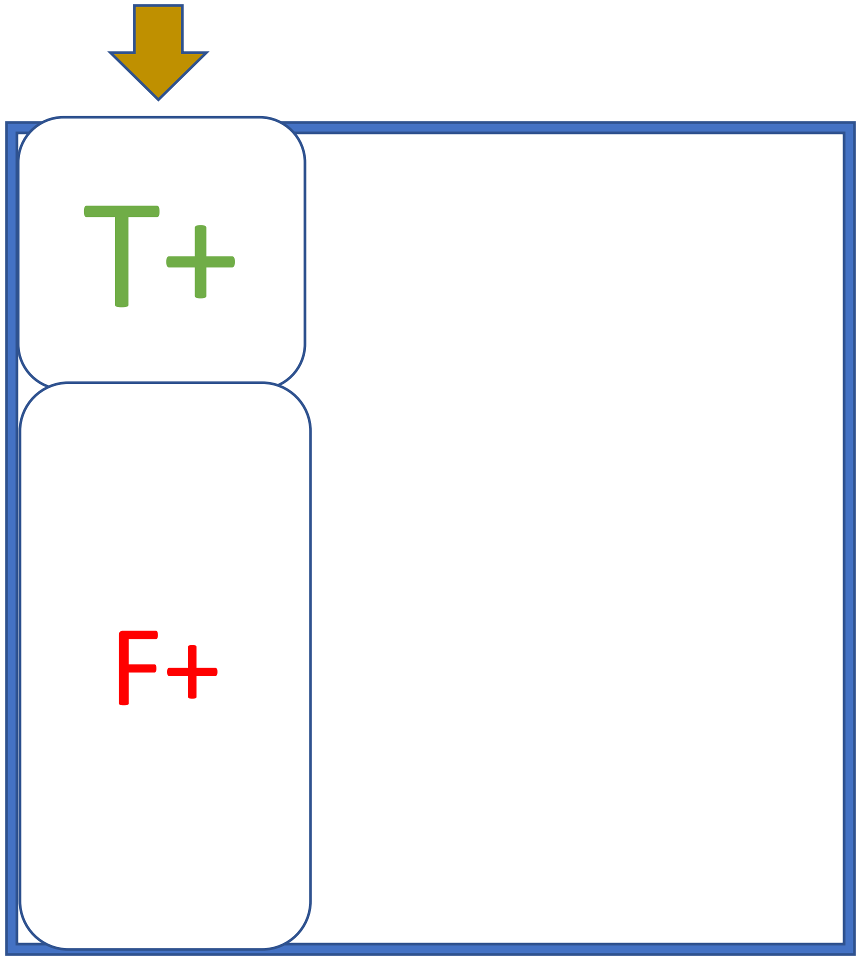

Focus on a single label

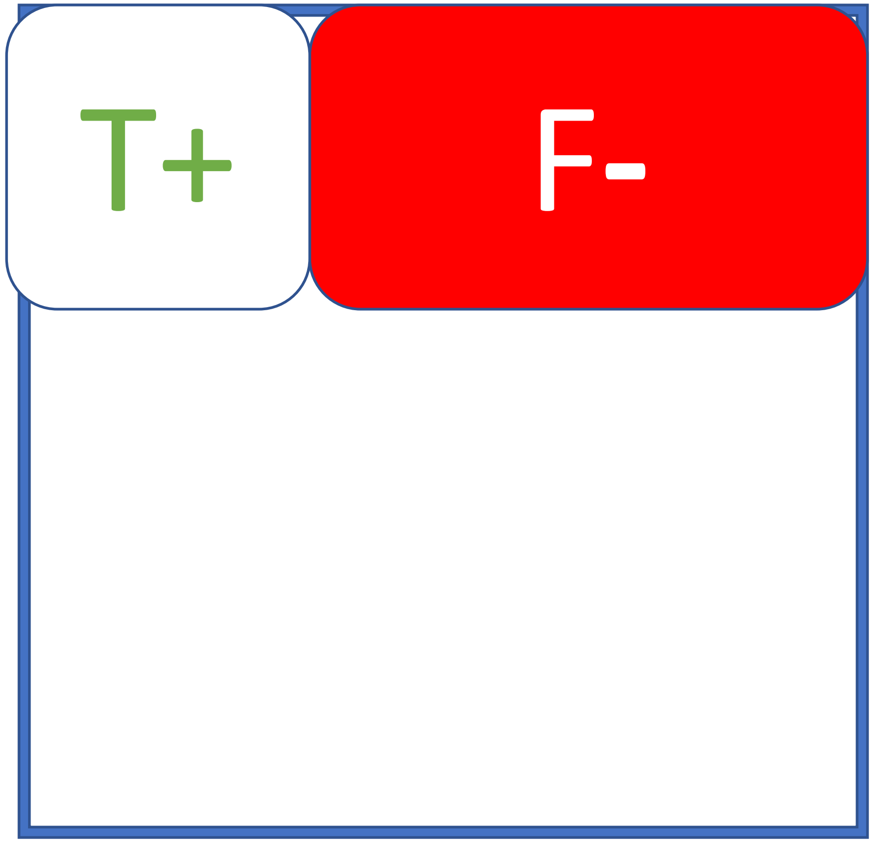

$$\textbf{Recall} = \frac{T+}{T+ \; + \; F-}$$

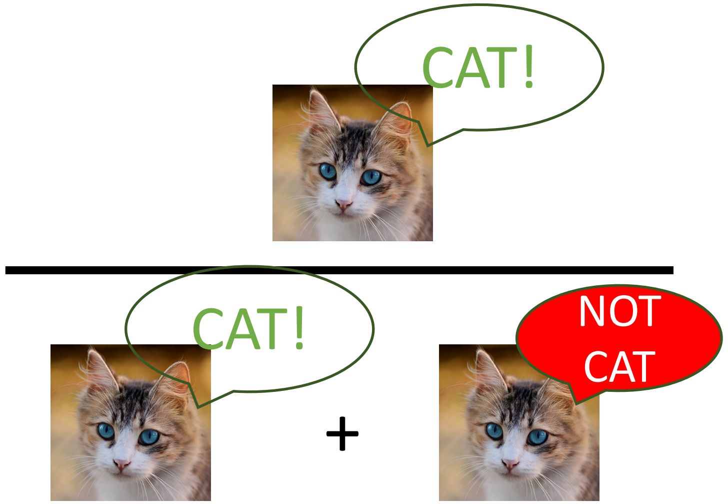

AKA sensitivity, hit rate

Focus on a single label

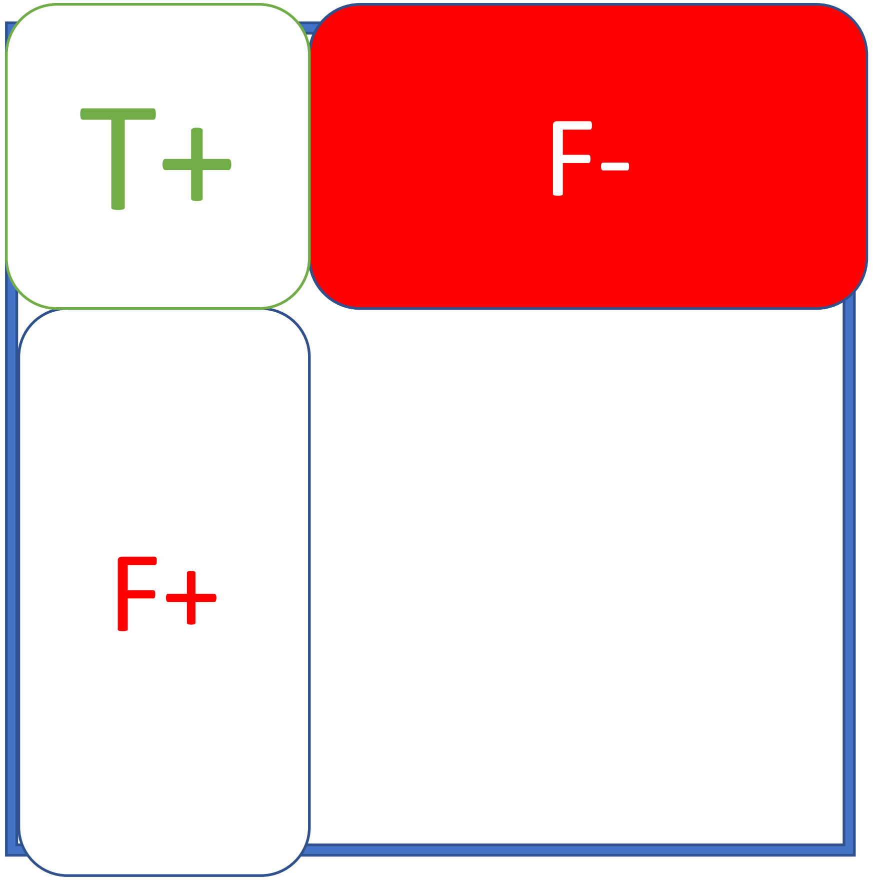

$$\textbf{Recall} = \frac{T+}{T+ \; + \; F-}$$

$$\textbf{Precision} = \frac{T+}{T+ \; + \; F+}$$

Focus on a single label

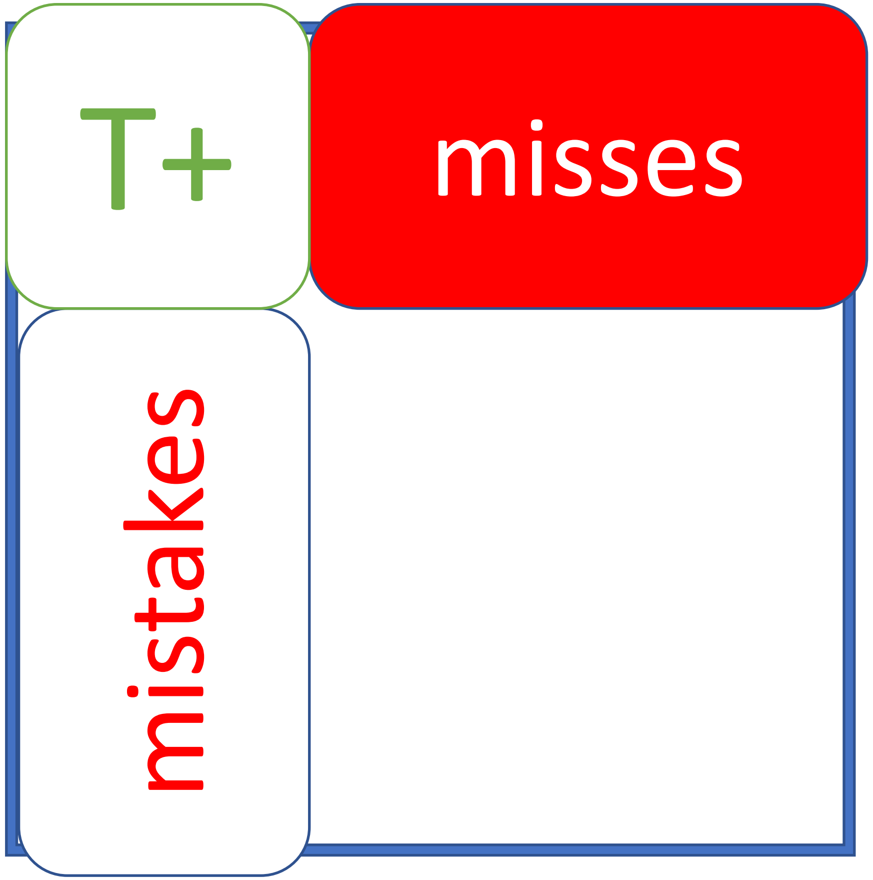

$$\text{Recall} = \frac{\text{T+}}{\text{T+} + \text{misses}}$$

$$\text{Precision} = \frac{\text{T+}}{\text{T+} + \text{mistakes}}$$

Focus on a single label

Hewlett Packard Enterprise Data Science Institute
UNIVERSITY OF HOUSTON

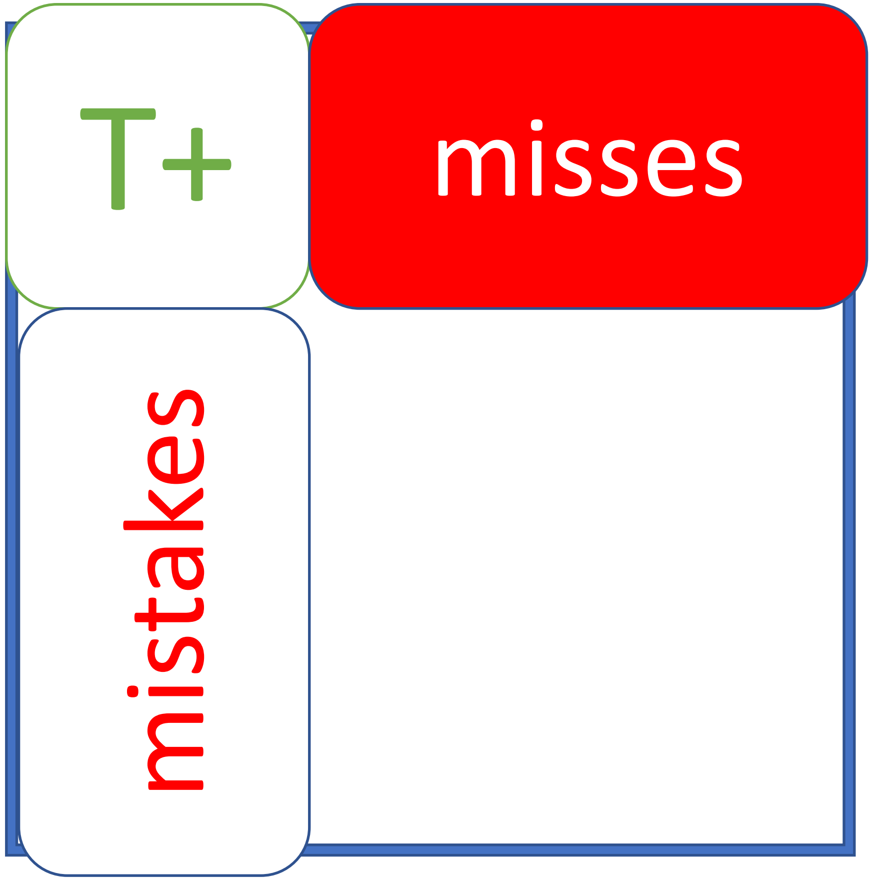

$$\textbf{Recall} = \frac{\boxed{T+}}{\boxed{T+} + \boxed{\text{misses}}}$$

$$F_1 = \frac{\boxed{T+}}{\boxed{T+} + \dfrac{\boxed{\text{mistakes}} + \boxed{\text{misses}}}{2}}$$

$$\textbf{Precision} = \frac{\boxed{T+}}{\boxed{T+} + \boxed{\text{mistakes}}}$$

Focus on a single label

$$F_1 = \frac{\boxed{T+}}{\boxed{T+} + \frac{\boxed{\text{mistakes}} + \boxed{\text{misses}}}{2}} = \frac{\boxed{T+}}{\frac{\boxed{T+} + \boxed{\text{misses}}}{2} + \frac{\boxed{T+} + \boxed{\text{mistakes}}}{2}}$$

$$= \frac{1}{\frac{1}{2}\left(\frac{1}{\text{recall}} + \frac{1}{\text{precision}}\right)}$$

$$= 2\frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Focus on a single label

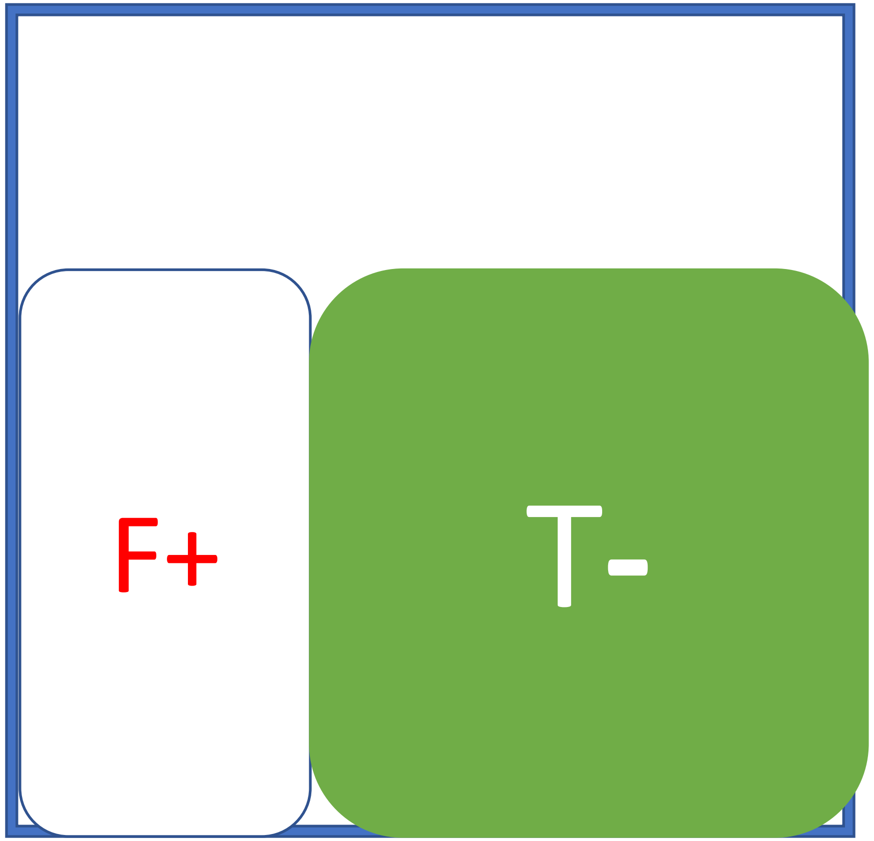

$$\textbf{Specificity} = \frac{T-}{F+ \; + \; T-}$$

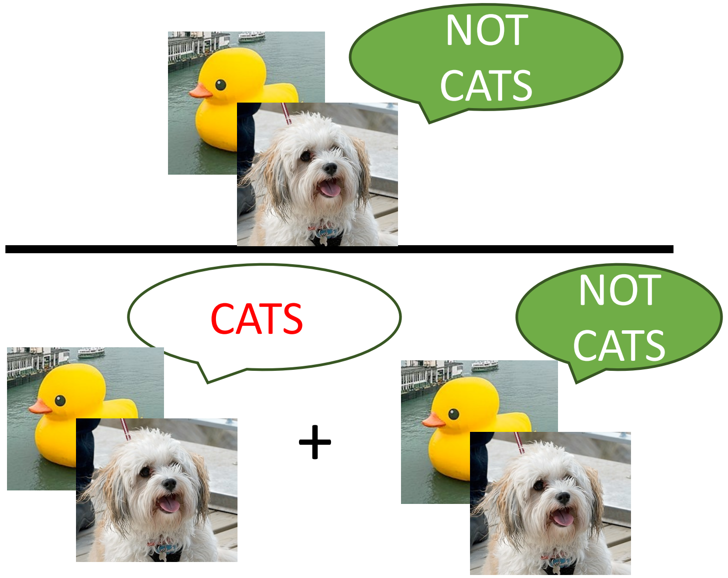

AKA selectivity

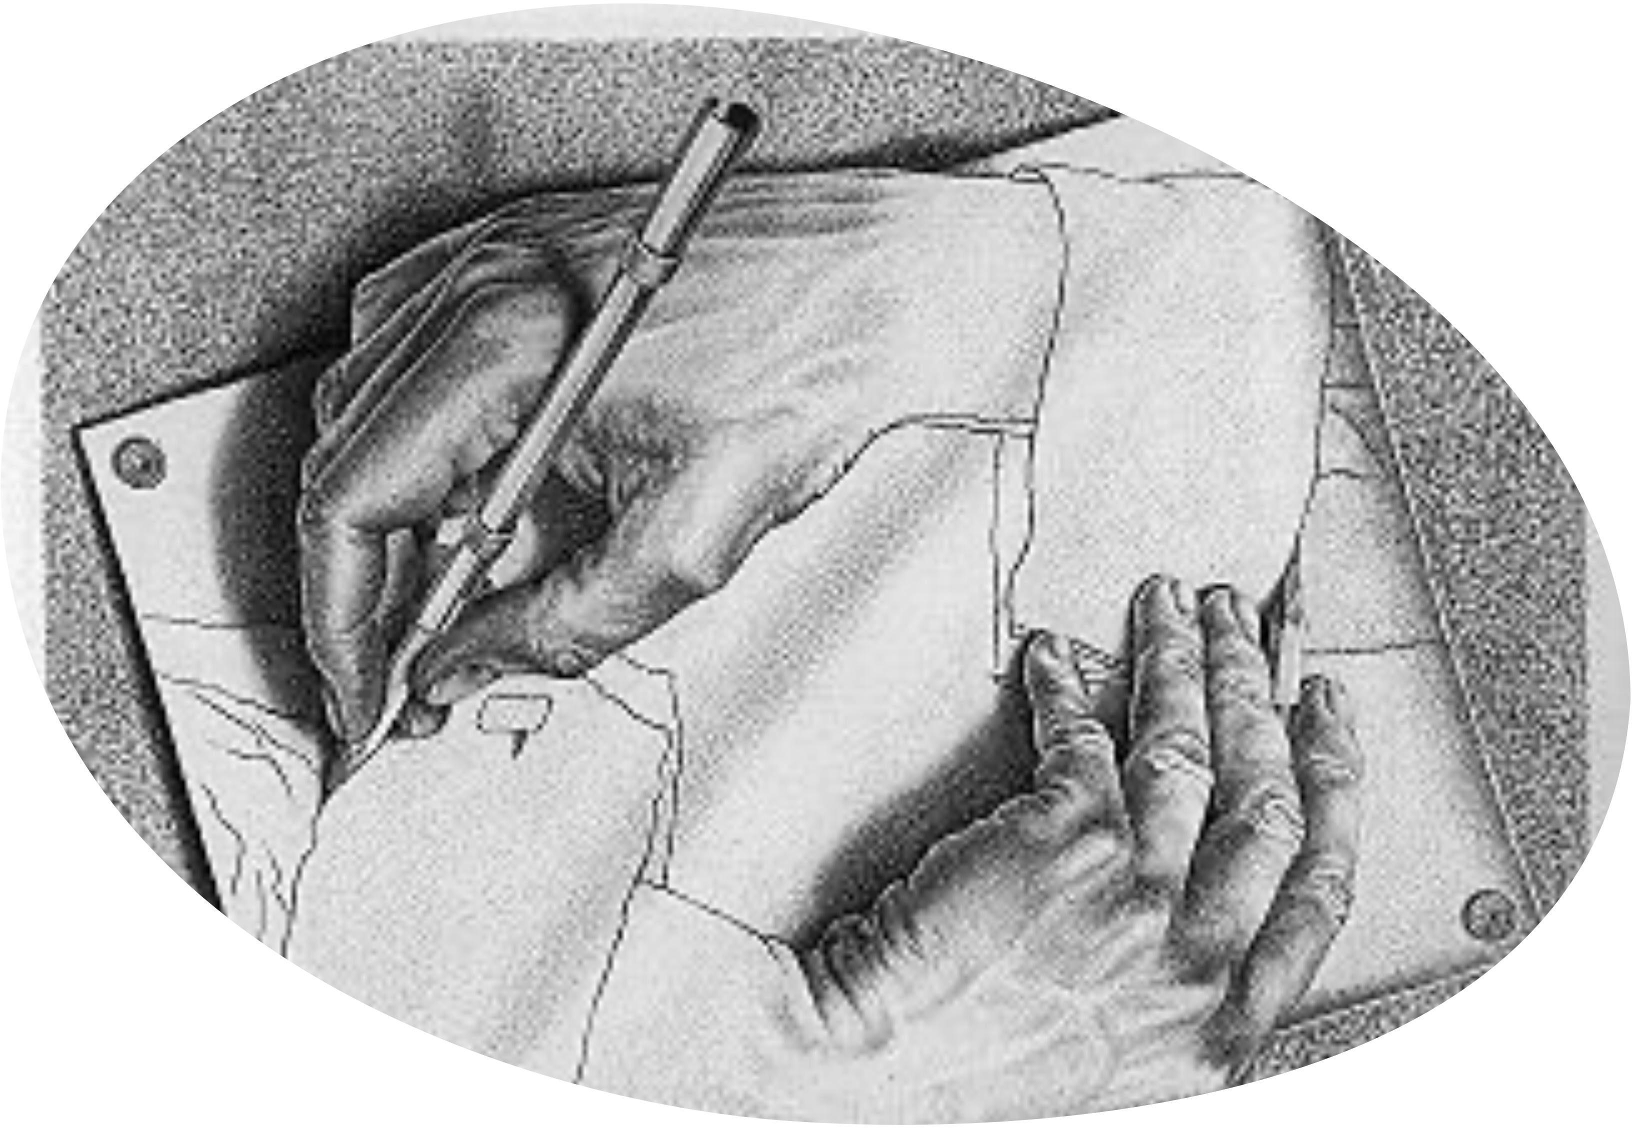

Hands-on Example:

Classification using k-NN + Logistic Regression

Confusion matrix

Plot_confusion_matrix(*estimator*, *X*, *y_true*,
labels=None,
sample_weight=None,
normalize=None,
display_labels=None,
include_values=True,
xticks_rotation='horizontal',
values_format=None,
cmap='viridis',
ax=None)

https://scikit-learn.org/stable/modules/model_evaluation.html#confusion-matrix

Confusion matrix

Labels: List of labels to index the matrix. This may be used to reorder or select a subset of labels. If None is given, those that appear at least once in y_true or y_pred are used in sorted order.

Normalize: Normalizes confusion matrix over the true (rows), predicted (columns) conditions or all the population. If None, confusion matrix will not be normalized.

include_values: Includes values in confusion matrix.

Classification Report

classification_report(*y_true*, *y_pred*,
labels=None,
target_names=None,
sample_weight=None,
digits=2,
output_dict=False,
zero_division='warn')

https://scikit-learn.org/stable/modules/generated/sklearn.metrics.classification_report.html

Classification Report

'macro’: Calculate metrics for each label, and find their unweighted mean. This does not take label imbalance into account.

'weighted’: Calculate metrics for each label, and find their average weighted by support (the number of true instances for each label). This alters ‘macro’ to account for label imbalance; it can result in an F-score that is not between precision and recall.

Note that if all labels are included, “micro”-averaging in a multiclass setting will produce precision and recall scores that are all identical to accuracy.

Hewlett Packard Enterprise Data Science Institute
UNIVERSITY OF HOUSTON

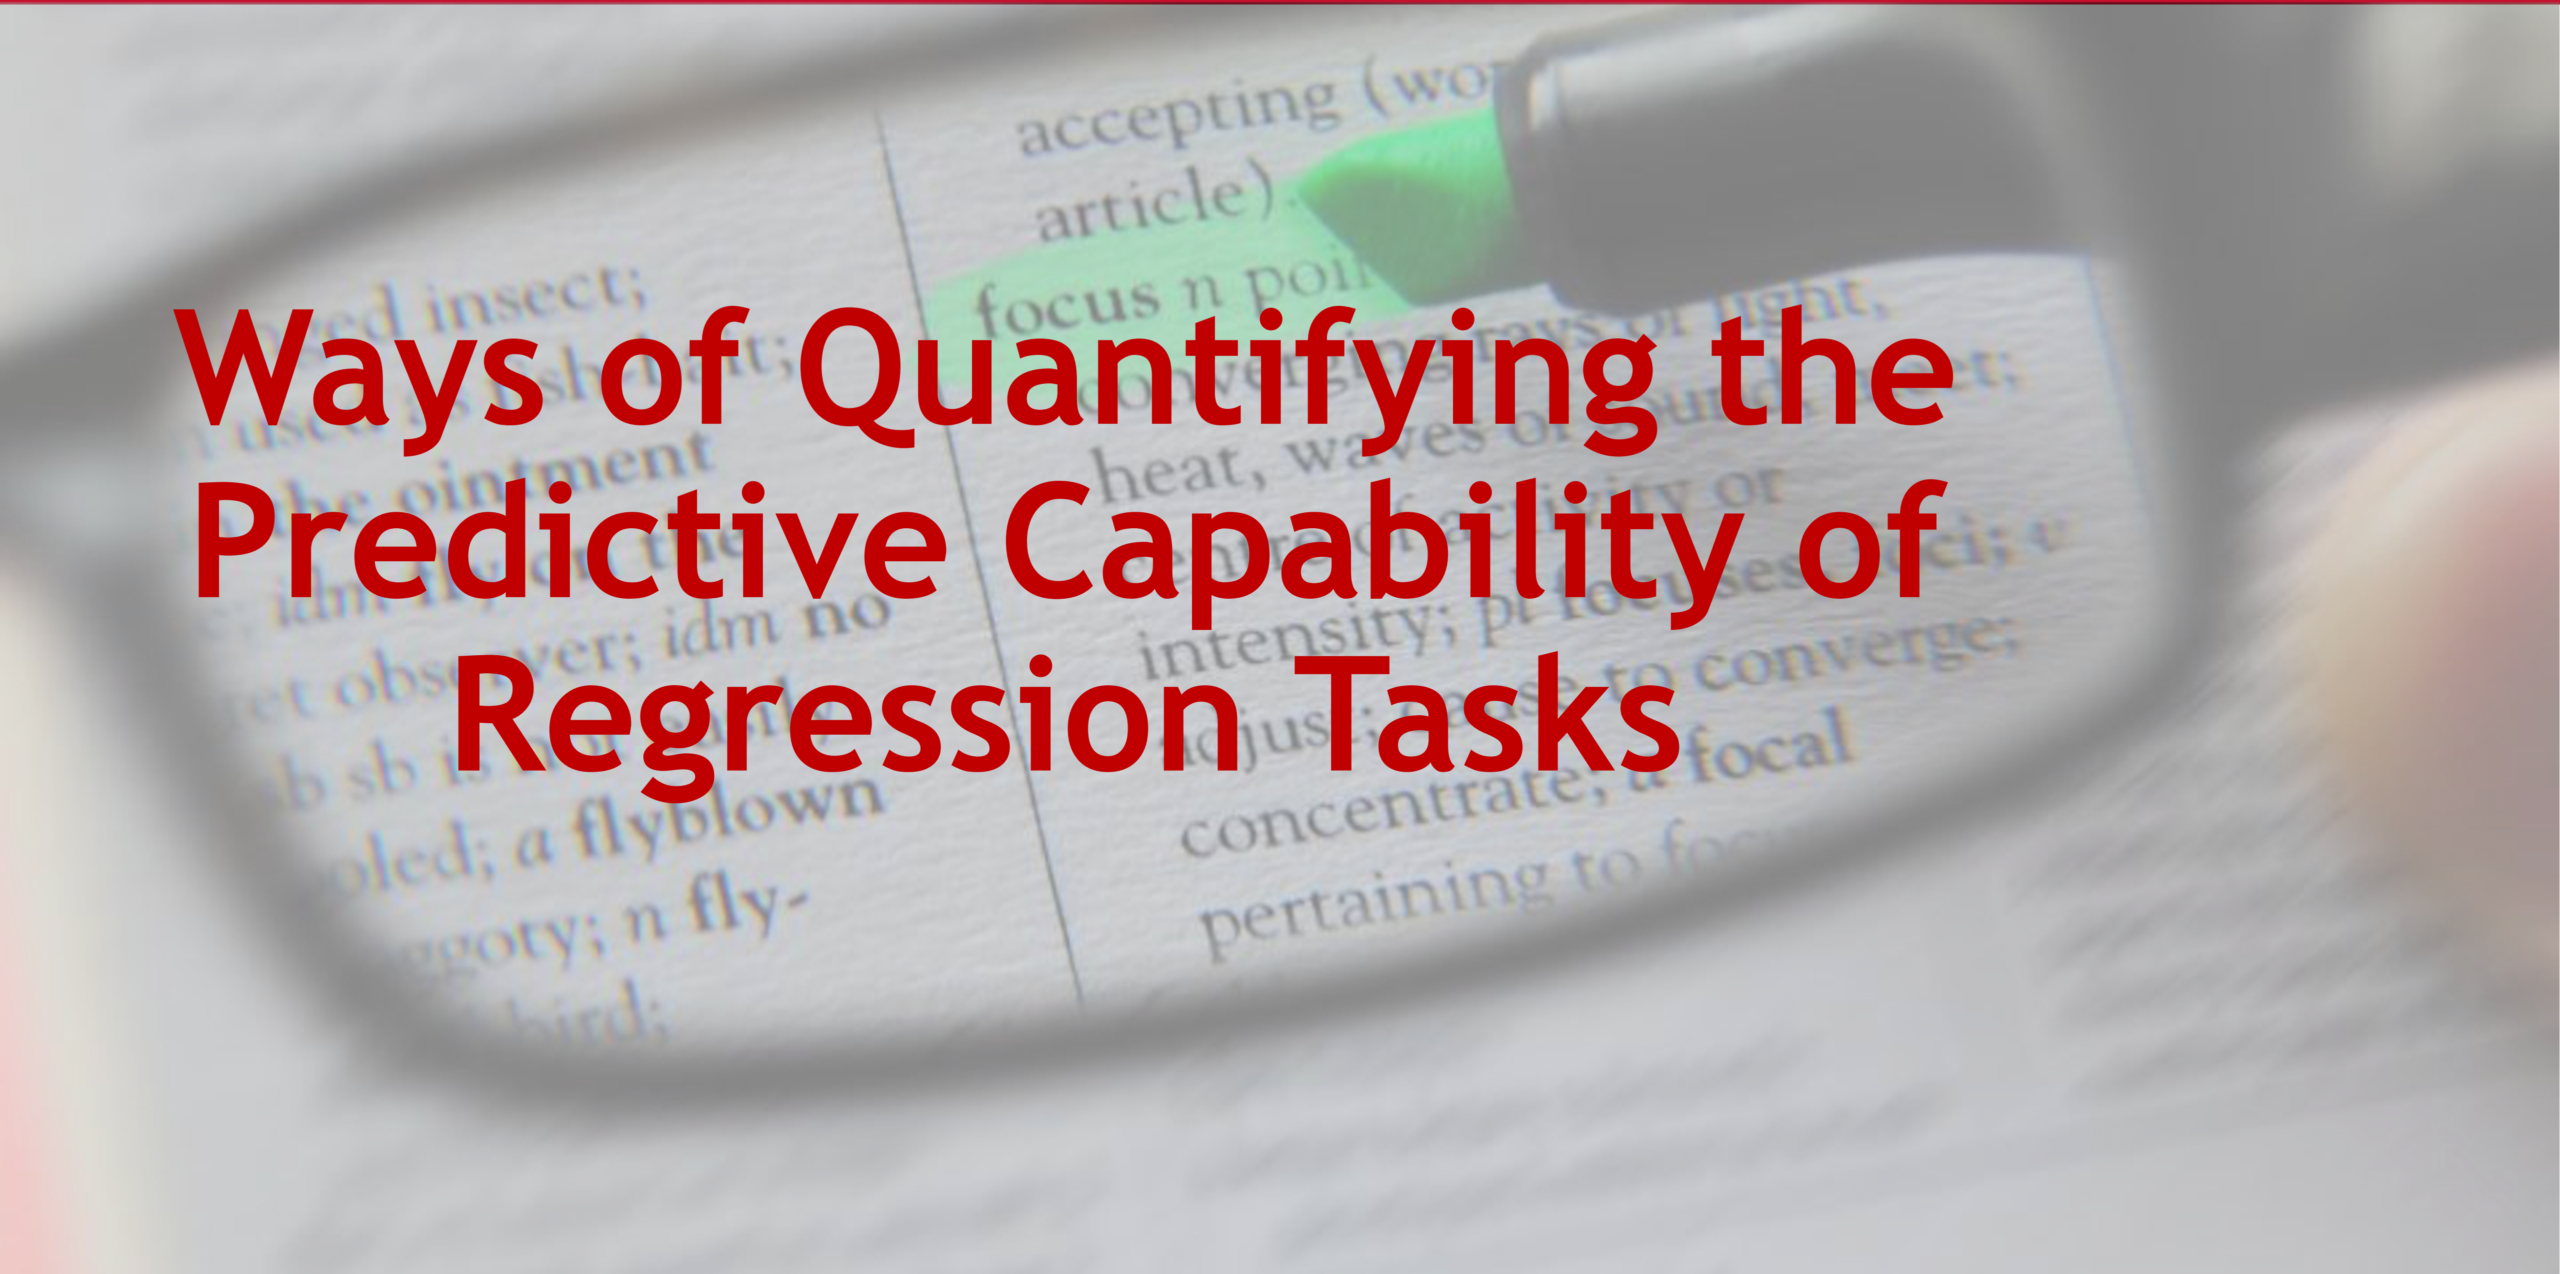

Ways of Quantifying the Predictive Capability of Regression Tasks

Mean Squared Error (MSE)

$$MSE = \frac{1}{N}\sum_{i=1}^{N}(f_i - y_i)^2$$

where N is the number of data points, f_i the value returned by the model and y_i the actual value for data point i.

Mean Squared Error (MSE)

$$MSE = \frac{1}{N}\boxed{\sum_{i=1}^{N}(f_i - y_i)^2}$$

where N is the number of data points, f_i the value returned by the model and y_i the actual value for data point i.

Euclidean distance squared, divided by number of points

Mean Squared Error (MSE)

$$MSE = \frac{1}{N}\boxed{\sum_{i=1}^{N}(f_i - y_i)^2}$$

where N is the number of data points, f_i the value returned by the model and y_i the actual value for data point i.

N=2

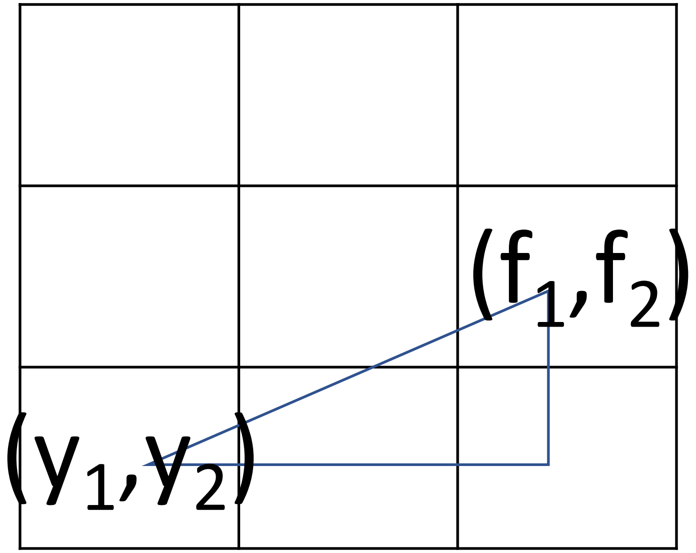

Mean Squared Error (MSE)

$$MSE = \frac{1}{N}\sum_{i=1}^{N}(f_i - y_i)^2$$

where N is the number of data points, f_i the value returned by the model and y_i the actual value for data point i.

N=2

2	2.236	2.828
1	1.414	2.236
(y_1, y_2)	1	2

Mean Absolute Deviation (MAD)

$$\frac{1}{N}\sum_{i=1}^{N}\left|f_i - y_i\right|$$

Mean Absolute Deviation (MAD)

Manhattan distance
divided by number of points

$$\frac{1}{N}\sum_{i=1}^{N}|f_i - y_i|$$

N=2

2	3	4
1	2	3
(y_1,y_2)	1	2

Maximum error

N=2

2	2	2
1	1	2
(y_1, y_2)	1	2

Remember the L^p norms (Minkowski)?

$$\|x\|_p = \left(|x_1|^p + |x_2|^p + \cdots + |x_n|^p\right)^{1/p}$$

$$\|x\|_\infty = \max\{|x_1|, |x_2|, \ldots, |x_n|\}$$

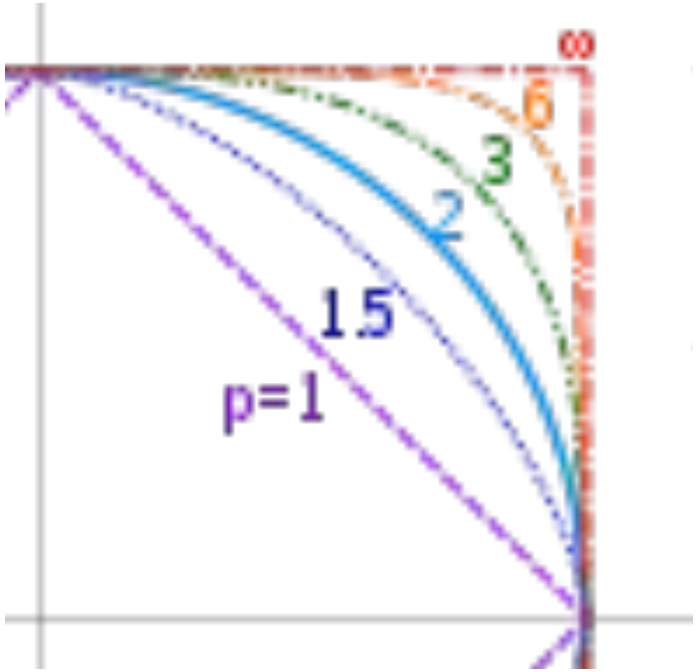

Unit circle for different values of p

Developed by Ioannis Konstantinidis

Hewlett Packard Enterprise Data Science Institute
UNIVERSITY OF HOUSTON

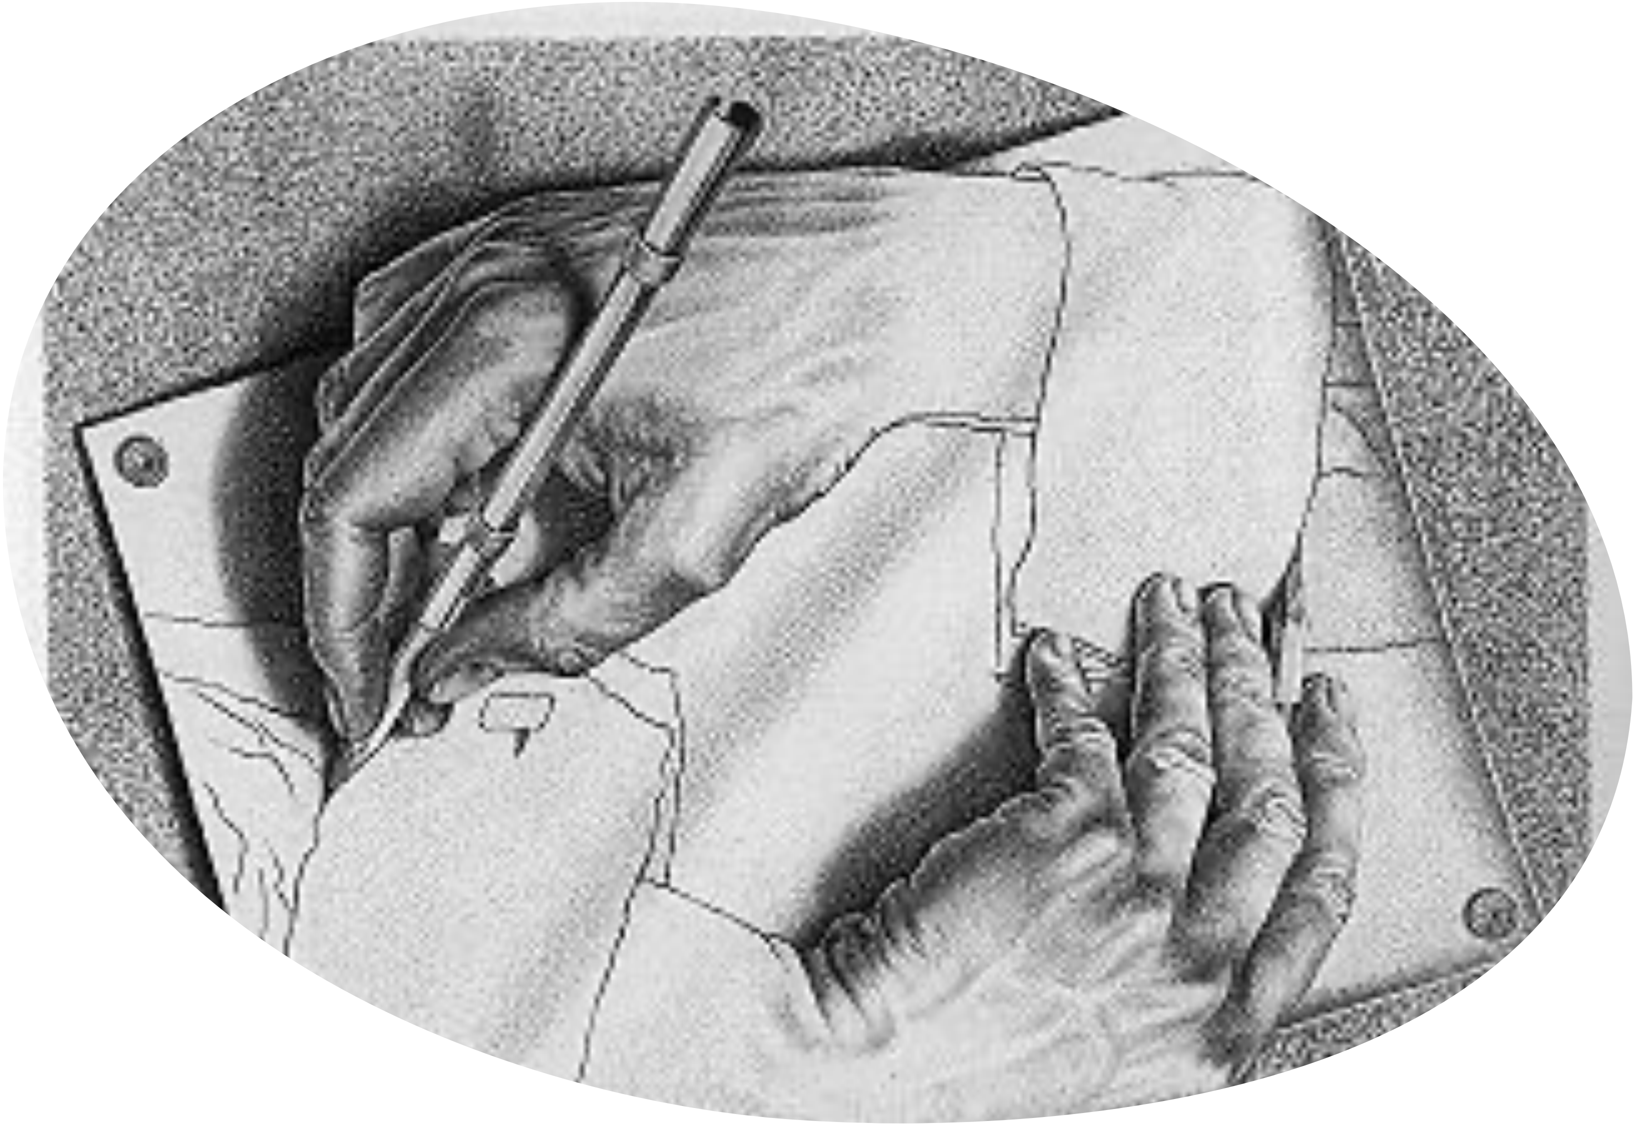

Hands-on Example:

Linear Regression

Homework Assignment #1
Due Monday, June 10, 11:59 pm (Central)

Regularization: mix and match

Mix and match

Multivariable Regression: F = X β + constant

$$F = A + B_1 X_1 + B_2 X_2 + \ \dots \ + \ B_K X_K$$

Mix and match

Multivariable Regression: F = X β + constant

$$\sum_{i=1}^{N}(f_i - y_i)^2$$

$$= (y - X\beta)^T(y - X\beta)$$

Mix and match

Multivariable Regression: F = X β + constant

$$\text{Ridge Cost} = (y - X\beta)^T(y - X\beta) + ||\beta||_2^2$$
$$\text{Lasso Cost} = (y - X\beta)^T(y - X\beta) + ||\beta||_1$$

Mix and match

Multivariable Regression: F = X β + constant

$$\text{Ridge Cost} = \underbrace{\overbrace{(y - X\beta)^T(y - X\beta)}^{L^2}}_{} + \overbrace{||\beta||_2^2}^{L^2}$$

$$\text{Lasso Cost} = \underbrace{(y - X\beta)^T(y - X\beta)}_{L^2} + \underbrace{||\beta||_1}_{L^1}$$

Mix and match

Multivariable Regression: F = X β + constant

$$\text{Ridge Cost} = \overbrace{(y - X\beta)^T(y - X\beta)}^{L^2} + \alpha \overbrace{||\beta||_2^2}^{L^2}$$

$$\text{Lasso Cost} = \underbrace{(y - X\beta)^T(y - X\beta)}_{L^2} + \alpha \underbrace{||\beta||_1}_{L^1}$$

Mix and match

Multivariable Regression: F = X β + constant

$$\text{Ridge Cost} = (y - X\beta)^T (y - X\beta) + \alpha ||\beta||_2^2$$
$$\text{Lasso Cost} = (y - X\beta)^T (y - X\beta) + \alpha ||\beta||_1$$

α is the regularization (hyper)parameter